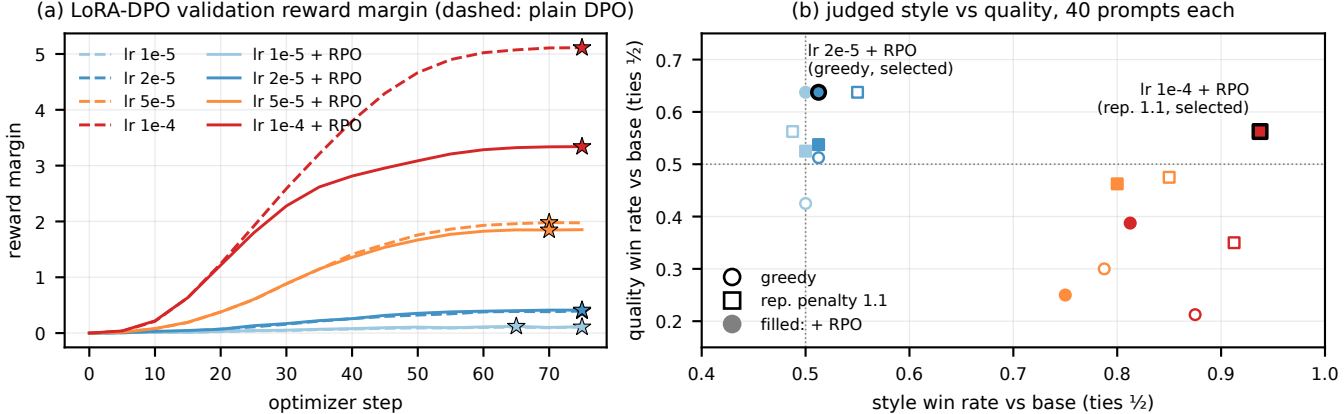
(a) LoRA-DPO validation reward margin (dashed: plain DPO)
lr 1e-5
lr 2e-5
lr 5e-5
lr 1e-4
lr 1e-5 + RPO
lr 2e-5 + RPO
lr 5e-5 + RPO
lr 1e-4 + RPO
reward margin
optimizer step
(b) judged style vs quality, 40 prompts each
lr 2e-5 + RPO
(greedy, selected)
lr 1e-4 + RPO
(rep. 1.1, selected)
greedy
rep. penalty 1.1
filled: + RPO
quality win rate vs base (ties ½)
style win rate vs base (ties ½)