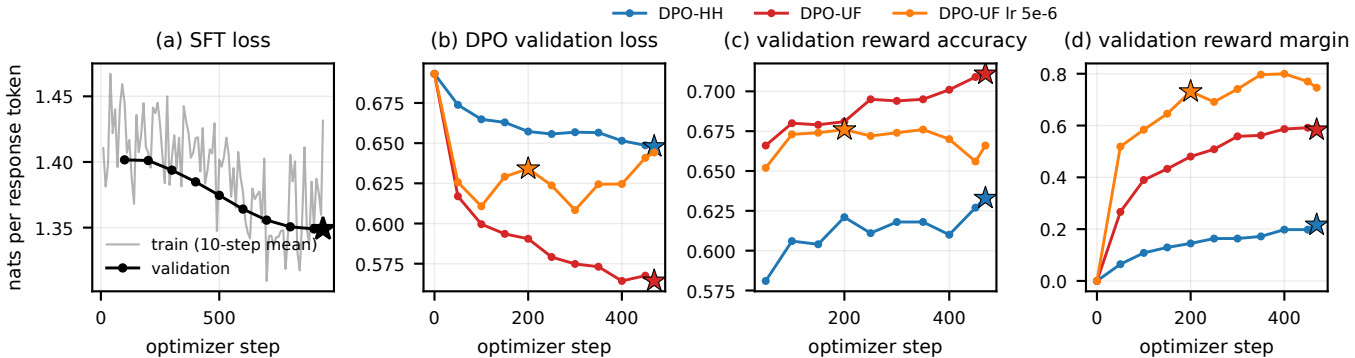
DPO-HH
DPO-UF
DPO-UF lr 5e-6
(a) SFT loss
(b) DPO validation loss
(c) validation reward accuracy
(d) validation reward margin
nats per response token
1.45
1.40
1.35
train (10-step mean)
validation
0
500
optimizer step
0.675
0.650
0.625
0.600
0.575
0
200
400
optimizer step
0.700
0.675
0.650
0.625
0.600
0.575
200
400
optimizer step
0.8
0.6
0.4
0.2
0.0
0
200
400
optimizer step